

```
In [1]: import numpy as np
import pandas as pd
pd.options.display.max_columns=100
from sklearn.model_selection import train_test_split, cross_val_score, cross_validate
import sklearn.metrics
from sklearn.preprocessing import StandardScaler as SSc
from sklearn.tree import DecisionTreeRegressor as DTR
from sklearn.ensemble import RandomForestRegressor as RFR
from sklearn.neighbors import KNeighborsRegressor as KNR
import matplotlib.pyplot as plt
import graphviz as gviz
%matplotlib inline

#set width of window to preference
from IPython.core.display import display, HTML
display(HTML("<style>.container { width:90% !important; }</style>"))
display(HTML("<style>.output.output_scroll{ height:100% !important; }</style>")) #breaks scroll output vertical so you see the whole output, disable this if you prefer.
```

```
In [2]: #year = "2019" #choose year to get data from
#split = "summer" #choose split to get data from (spring, summer, worlds)
#infile = r"C:\Users\Triplea657\000 MSCS-335 2020\Datasets\League_" #path
#inf = "-Wrangled.csv" #file to read
#filein = infile+year+"\\"+year+'-'+split+'-'+inf
#data = pd.read_csv(filein, low_memory=False)
#data.head(10)

#changed for submission version
data = pd.read_csv("Datasets/League_2019/2019-summer-Wrangled.csv", index_col=0, low_memory=False)
data.head()
```

Out[2]:

	league_CBLol	league_LCK	league_LCS	league_LEC	league_LMS	gamelength	result	k
0	0.0	0.0	1.0	0.0	0.0	35.500000	1.0	21.0
1	0.0	0.0	1.0	0.0	0.0	35.500000	0.0	14.0
2	0.0	0.0	1.0	0.0	0.0	29.700000	1.0	11.0
3	0.0	0.0	1.0	0.0	0.0	29.700000	0.0	4.0
4	0.0	0.0	1.0	0.0	0.0	31.983333	1.0	12.0

```
In [3]: var = []
        for i in data:
            var.append(i)
        print(var)
        for i, v in enumerate(var):
            print(i, v)
```

```
['league_CBLol', 'league_LCK', 'league_LCS', 'league_LEC', 'league_LMS',  
'gamelength', 'result', 'k', 'd', 'a', 'fb', 'kpm', 'okpm', 'ckpm',  
'fd', 'fdtime', 'teamdragkills', 'oppdragkills', 'elementals', 'oppelem  
entals', 'firedrakes', 'waterdrakes', 'earthdrakes', 'airdrakes', 'elde  
rs', 'oppelders', 'herald', 'heraldtime', 'ft', 'fttime', 'firstmidoute  
r', 'firstttothreetowers', 'teambaronkills', 'oppbaronkills', 'dmgtocham  
ps', 'dmgtochampsperminute', 'wards', 'wpm', 'wardkills', 'wcpm', 'tota  
lgold', 'earnedgpm', 'goldspent', 'gspd', 'monsterkillsownjungle', 'mon  
sterkillsenemyjungle', 'cspm', 'goldat10', 'oppgoldat10', 'gdat10', 'go  
ldat15', 'oppgoldat15', 'gdat15', 'xpat10', 'oppxpat10', 'xpdatt10', 'cs  
at10', 'oppcsat10', 'csdat10', 'csat15', 'oppcsat15', 'csdat15']  
0 league_CBLol  
1 league_LCK  
2 league_LCS  
3 league_LEC  
4 league_LMS  
5 gamelength  
6 result  
7 k  
8 d  
9 a  
10 fb  
11 kpm  
12 okpm  
13 ckpm  
14 fd  
15 fdtime  
16 teamdragkills  
17 oppdragkills  
18 elementals  
19 oppelementals  
20 firedrakes  
21 waterdrakes  
22 earthdrakes  
23 airdrakes  
24 elders  
25 oppelders  
26 herald  
27 heraldtime  
28 ft  
29 fttime  
30 firstmidouter  
31 firstttothreetowers  
32 teambaronkills  
33 oppbaronkills  
34 dmgtochamps  
35 dmgtochampsperminute  
36 wards  
37 wpm  
38 wardkills  
39 wcpm  
40 totalgold  
41 earnedgpm  
42 goldspent  
43 gspd  
44 monsterkillsownjungle  
45 monsterkillsenemyjungle  
46 cspm  
47 goldat10  
48 oppgoldat10  
49 gdat10  
50 goldat15  
51 oppgoldat15  
52 gdat15
```

1 - Trying to predict the length of the game played

Split data into X,Y where Y is the length of the game, then normalize the inputs.

```
In [4]: X = data.iloc[:,data.columns != 'gamelength']
Y = data.iloc[:,data.columns == 'gamelength']
#transform input data (normalize scaling)
ssc = Ssc()
Xft = ssc.fit_transform(X)
X = pd.DataFrame(Xft)
print("Xtr(Xtrain),Xtst(Xtest),Ytr(Ytrain),Ytst(Ytest) shapes: ")
Xtr,Xtst,Ytr,Ytst = train_test_split(X,Y.values.ravel(),test_size=0.2,ran
dom_state=2020)
print(Xtr.shape,Xtst.shape,Ytr.shape,Ytst.shape)

Xtr(Xtrain),Xtst(Xtest),Ytr(Ytrain),Ytst(Ytest) shapes:
(1155, 61) (289, 61) (1155,) (289,)
```

```
In [5]: print("Classifier scores:\n"+"-"*18+"\n")

tree = DTR(max_depth=5)
tree.fit(Xtr,Ytr)
scr = cross_val_score(tree,Xtst,Ytst, cv=5)
print("Tree \nscore avg:"+str(sum(scr)/5)+"\nscore = "+str(scr)+"\n\
n"+"-"*64)

for i in range(1,20,2):
    forest = RFR(n_estimators=i,max_depth=5)
    forest.fit(Xtr,Ytr)
    scr = cross_val_score(forest,Xtst,Ytst, cv=5)
    print("\nRandom Forest trees = "+str(i)+" depth = 5 \nscore avg: "+str(
r(sum(scr)/5)+" \nscores: "+str(scr))

print("\n"+"-"*64)

for i in range(1,16,3):
    knn = KNR(n_neighbors=i)
    knn.fit(Xtr,Ytr)
    scr = cross_val_score(knn,Xtst,Ytst, cv=5)
    print("\nK-Nearest Neighbors "+str(i)+"-neighbors\nscore avg:"+str(su
m(scr)/5)+"\nscore = "+
        str(scr))
```

Classifier scores:

Tree

score avg:0.8688339362334357

score = [0.90064553 0.84418264 0.89714237 0.89148452 0.81071462]

Random Forest trees = 1 depth = 5

score avg: 0.8164661368933335

scores: [0.88549141 0.82110831 0.82165248 0.83091341 0.72316508]

Random Forest trees = 3 depth = 5

score avg: 0.9066697193082612

scores: [0.91099621 0.91309417 0.90988224 0.91149418 0.88788179]

Random Forest trees = 5 depth = 5

score avg: 0.894486939135383

scores: [0.94675132 0.89394793 0.8601876 0.88808896 0.88345888]

Random Forest trees = 7 depth = 5

score avg: 0.9155150617489675

scores: [0.90452363 0.92936666 0.90838157 0.91841461 0.91688883]

Random Forest trees = 9 depth = 5

score avg: 0.9167504015280432

scores: [0.95102155 0.93116634 0.88536546 0.92816278 0.88803588]

Random Forest trees = 11 depth = 5

score avg: 0.9187686643099913

scores: [0.93674206 0.92731609 0.90215989 0.9298727 0.89775258]

Random Forest trees = 13 depth = 5

score avg: 0.9263979910877467

scores: [0.94186259 0.93119973 0.90635365 0.92557475 0.92699923]

Random Forest trees = 15 depth = 5

score avg: 0.9273671560286475

scores: [0.93883081 0.93474997 0.91045486 0.92222491 0.93057523]

Random Forest trees = 17 depth = 5

score avg: 0.9272029254562095

scores: [0.94320471 0.93716017 0.90971531 0.92574404 0.9201904]

Random Forest trees = 19 depth = 5

score avg: 0.9285613689890869

scores: [0.9397332 0.93297646 0.91604547 0.92532215 0.92872957]

K-Nearest Neighbors 1-neighbors

score avg:0.5503541585117673

score = [0.43000924 0.71837904 0.50965417 0.59119001 0.50253834]

K-Nearest Neighbors 4-neighbors

score avg:0.7242329981047385

score = [0.68671451 0.79204271 0.76680979 0.71530448 0.6602935]

K-Nearest Neighbors 7-neighbors

score avg:0.7094477063937197

score = [0.67021646 0.74731298 0.74034743 0.68501081 0.70435085]

K-Nearest Neighbors 10-neighbors

K-Nearest Neighbors performs quite poorly, but trees perform fairly well and random forests perform very well even with fairly small numbers of trees.

2 - Trying to predict the number of kills, deaths and assists for the current team in the current game

Split data into X,Y where Y is the kills, deaths and assists, then normalize the inputs.

```
In [6]: idx = [7,8,9]
idxtitles = ['kills', 'deaths', 'assists']
X = []
Y = []
for i in idx:
    X.append(data.drop(data.columns[idx],axis=1))
    Y.append(data.iloc[:,i])
    #transform input data (normalize scaling)
    ssc = SSc()
    Xft = ssc.fit_transform(X[i-7])
    X[i-7] = pd.DataFrame(Xft)
    print("Xtr(Xtrain),Xtst(Xtest),Ytr(Ytrain),Ytst(Ytest) shapes: ")
    Xtr,Xtst,Ytr,Ytst = train_test_split(X[i-7],Y[i-7],test_size=0.2,random_state=2020)
    print(Xtr.shape,Xtst.shape,Ytr.shape,Ytst.shape)

Xtr(Xtrain),Xtst(Xtest),Ytr(Ytrain),Ytst(Ytest) shapes:
(1155, 59) (289, 59) (1155,) (289,)
Xtr(Xtrain),Xtst(Xtest),Ytr(Ytrain),Ytst(Ytest) shapes:
(1155, 59) (289, 59) (1155,) (289,)
Xtr(Xtrain),Xtst(Xtest),Ytr(Ytrain),Ytst(Ytest) shapes:
(1155, 59) (289, 59) (1155,) (289,)
```

```

In [7]: #
-----

#---begin code
#
-----

breakline = "-"*64

scoring = {'FVE': 'explained_variance',
           'MSE': 'neg_mean_squared_error',
           'R2': 'r2'}
def cscore(model,X,Y):
    cr_v = cross_validate(model, X, Y, scoring=scoring,cv=5, return_train
_score=False)
    return cr_v
def tst(X,Y):
    for i in range(2):
        dpth = 3+i
        tree = DTR(max_depth=dpth)
        tree.fit(Xtr,Ytr)
        scr = cscore(tree,Xtst,Ytst)
        print("\nThe optimal depth is 3-4 for a single tree.\n"+"-"*10+"T
ree of depth {} scores:".format(dpth))
        for j,k in enumerate(scr.keys()):
            if j > 1:
                if(k=='test_MSE'):
                    print("-----{} (0.0 is best)\nscores:    {}\navg scor
e: {}".format(k,scr[k],scr[k].mean()))
                else:
                    print("-----{} (1.0 is best)\nscores:    {}\navg scor
e: {}".format(k,scr[k],scr[k].mean()))

        print("\n\n"+"-"*36)

        print("\nK-Nearest Neighbors reaches near maximum accuracy at 5 neigh
bors and rising marginally until 9 neighbors")
        for l in range(3,12,2):
            knn = KNR(n_neighbors=l)
            knn.fit(Xtr,Ytr)
            scr = cscore(knn,Xtst,Ytst)
            print("\n"+"-"*10+"K-Nearest Neighbors, {}-neighbors scores:".for
mat(l))
            for j,k in enumerate(scr.keys()):
                if j > 1:
                    if(k=='test_MSE'):
                        print("-----{} (0.0 is best)\nscores:    {}\navg scor
e: {}".format(k,scr[k],scr[k].mean()))
                    else:
                        print("-----{} (1.0 is best)\nscores:    {}\navg scor
e: {}".format(k,scr[k],scr[k].mean()))

            print("\n\n"+"-"*36)

            dpth=4
            print("\nRandom forests seem to reach a maximum accuracy of about 0.8
4 for FVE and R^2 with any parameters.\nIt reaches optimal accuracy most
efficiently at ~10 trees of depth 3 and ~5 trees of depth 4.")
            for l in range(1,16,5): #change number of trees
                forest = RFR(n_estimators=l,max_depth=dpth)

```



```
I used a Decision Tree, K-Nearest Neighbors, and Random Forests to predict ['kills', 'deaths', 'assists']
```

```
FVE best score: 1.0
MSE best score: 0.0, negative indicates that 0.0 is the best score as opposed to 1.0
FVE best score: 1.0
```

```
-----
kills regressor scores:
-----
```

```
The optimal depth is 3-4 for a single tree.
-----Tree of depth 3 scores:
-----test_FVE (1.0 is best)
scores: [0.74296769 0.82130487 0.83242113 0.70568816 0.73976042]
avg score: 0.7684284535774695
-----test_MSE (0.0 is best)
scores: [-57.98651027 -45.50304525 -39.9737957 -43.8362467 -64.07591737]
avg score: -50.27510305564451
-----test_R2 (1.0 is best)
scores: [0.71647788 0.81876985 0.83023975 0.70354059 0.72240388]
avg score: 0.7582863895953074
```

```
The optimal depth is 3-4 for a single tree.
-----Tree of depth 4 scores:
-----test_FVE (1.0 is best)
scores: [0.77095832 0.82403861 0.75857904 0.67976189 0.77320399]
avg score: 0.7613083688845369
-----test_MSE (0.0 is best)
scores: [-54.02552957 -44.2356364 -57.58504619 -47.55941503 -54.91570954]
avg score: -51.66426734736389
-----test_R2 (1.0 is best)
scores: [0.7358449 0.8238177 0.75544849 0.67836124 0.76208865]
avg score: 0.7511121965432139
```

```
-----
K-Nearest Neighbors reaches near maximum accuracy at 5 neighbors and rising marginally until 9 neighbors
```

```
-----K-Nearest Neighbors, 3-neighbors scores:
-----test_FVE (1.0 is best)
scores: [0.67132017 0.5897223 0.542423 0.50447341 0.42788093]
avg score: 0.5471639600628959
-----test_MSE (0.0 is best)
scores: [-69.27011494 -104.1302682 -108.79310345 -73.36781609 -142.06822612]
avg score: -99.52590576056998
-----test_R2 (1.0 is best)
scores: [0.66130726 0.58526854 0.53797871 0.50382205 0.38451777]
avg score: 0.5345788656771712
```

```
-----K-Nearest Neighbors, 5-neighbors scores:
-----test_FVE (1.0 is best)
scores: [0.71106917 0.58881426 0.65815358 0.55178016 0.46796877]
avg score: 0.5955571874020688
-----test_MSE (0.0 is best)
scores: [-61.32206897 -103.30689655 -80.54551724 -66.42827586 -134.19929825]
```

Much like the previous problem, K-nearest neighbors performed by far the worse and random forests performed very well even at low tree counts.